

Penn OS Companion Document: Group 19

NOTE: PENNFAT IS AT THE END

NOTE: NEW FILE DEFINED ON EACH NEW PAGE

pennos.c

void setTimerAlarmHandler(void)

Input Arguments: nothing

Return Values: No return values

int main(int argc, char argv)**

Input Arguments: number of arguments, array of arguments in strings

Return Values: returns 0

dependencies.h

#pragma once

#include <stdio.h>

// Define the structure for a Process

typedef struct process{

struct pcb* pcb;

struct process* next;

} Process;

Declare the heads and tails of all the queues being used (high, medium, low, blocked, stopped, zombie, orphan, temp

extern Process *highQhead;

extern Process *highQtail;

extern Process *medQhead;

extern Process *medQtail;

extern Process *lowQhead;

extern Process *lowQtail;

extern Process *blockedQhead;

extern Process *blockedQtail;

extern Process *stoppedQhead;

extern Process *stoppedQtail;

extern Process *zombieQhead;

extern Process *zombieQtail;

extern Process *orphanQhead;

extern Process *orphanQtail;

extern Process *tempHead;

extern Process *tempTail;

extern int ticks;

extern int fgpid;

static const int quantum = 100000;

kernel.h

#pragma once

#include <ucontext.h> // getcontext, makecontext, setcontext, swapcontext

#include "pcb.h"

#include "scheduler.h"

#include "shell.h"

#include "user.h"

#include "dependencies.h"

#define S_SIGTERM 1

#define S_SIGSTOP 2

#define S_SIGCONT 3

struct pcb* k_process_create(struct pcb *parent);

Input Arguments: PCB struct of the parent process

Return Values: Returns the PCB of the child that was created

void k_process_cleanup(Process *p);

Input Arguments: Process p that is to be cleaned up

Return Values: No return values

int k_process_kill(Process *p, int signal);

Input Arguments: Process p to be killed, signal value of signal

Return Values: returns success or failure value

Process *findProcessByPid(int pid);

Input Arguments: pid of process to find

Return Values: returns process that has been found

pcb

#pragma once

#include <ucontext.h> // getcontext, makecontext, setcontext, swapcontext

#include <stdio.h>

#include <stdlib.h>

#include "user.h"

#include "../fs/user.h"

Status values of the various processes

#define ZOMBIE 5

#define BLOCKED 4

#define STOPPED 3

#define RUNNING 2

#define SIG_TERMINATED 1

#define TERMINATED 0

Foreground and Background definitions

#define FG 0

#define BG 1

#define MAX_FILES 512

extern int pidCounter; // keeps track of pid to give to new processes

struct pcb {

ucontext_t context; // stores the context

int jobID; // stores the pennshell job ID

int numChild; // stores the number of children

int pid; // stores the pid

int ppid; // stores the parent pid

int waitChild; // stores pid of the child being waited on currently

int priority; // stores the priority of the process as an integer (0, 1, -1)

char *argument; // stores the argument of the process

int status; // stores the current status of the process

```
int bgFlag; // stores the background or foreground status of the process
int *childPids; // list of all pids in the job
int *childPidsFinished; // boolean array list that checks every pid is finished
int sleep_time_remaining; // number of ticks left for sleep to terminate gracefully
int changedStatus; // whether or not the process has changed status after creation or last wait
file_t *fd_table[MAX_FILES]; // fd table
};
```

```
char* strCopy(char* src, char* dest);
```

Input Arguments: source string, destination string

Return Values: destination string where the original one has been copied

```
struct pcb *initPCB();
```

Input Arguments: Nothing

Return Values: Struct pcb of the penn shell

```
struct pcb *createPcb(ucontext_t context, int pid, int ppid, int priority, int status);
```

Input Arguments: Context of the process, pid of the process, parents pid, priority of process, status process should start in

Return Values: Struct pcb of the new process

```
void freePcb(struct pcb *pcb_obj);
```

Input Arguments: pcb of the target process

Return Values: nothing

scheduler.h

```
#pragma once

#include <signal.h>
#include <stdio.h>
#include <stdbool.h>
#include <stdlib.h>
#include <sys/time.h>
#include <ucontext.h>
#include <unistd.h>
#include <valgrind/valgrind.h>
#include "pcb.h"
#include "kernel.h"
#include "user.h"
#include "shell.h"
#include "user_functions.h"
#include "dependencies.h"
```

// high, medium and low priority definitions

```
#define PRIORITY_HIGH -1
#define PRIORITY_MED 0
#define PRIORITY_LOW 1
```

```
extern ucontext_t schedulerContext; // context to return to scheduler
extern ucontext_t *activeContext; // context to run the active process
extern ucontext_t idleContext; // context to run the idle process
extern ucontext_t terminateContext; // context to run the terminate process function
```

```
void terminateProcess(void);
```

Input Arguments: nothing

Return Values: nothing

void scheduler(void);

Input Arguments: nothing

Return Values: nothing

void initContext(void);

Input Arguments: nothing

Return Values: nothing

void enqueueBlocked(Process* newProcess);

Input Arguments: Process struct of process to be enqueued

Return Values: nothing

void enqueueStopped(Process* newProcess);

Input Arguments: Process struct of process to be enqueued

Return Values: nothing

void enqueue(Process* newProcess);

Input Arguments: Process struct of process to be enqueued

Return Values: nothing

void enqueueZombie(Process* newProcess);

Input Arguments: Process struct of process to be enqueued

Return Values: nothing

void dequeueZombie(Process* newProcess);

Input Arguments: Process struct of process to be dequeued

Return Values: nothing

void dequeueBlocked(Process* newProcess);

Input Arguments: Process struct of process to be dequeued

Return Values: nothing

`void dequeueStopped(Process* newProcess);`

Input Arguments: Process struct of process to be dequeued

Return Values: nothing

`void dequeue(Process* newProcess);`

Input Arguments: Process struct of process to be dequeued

Return Values: nothing

`void iterateQueue(Process *head);`

Input Arguments: Process struct of head of the queue to be iterated through

Return Values: nothing

`void alarmHandler();`

Input Arguments: nothing

Return Values: nothing

`void setTimer(void);`

Input Arguments: nothing

Return Values: nothing

`void freeStacks(struct pcb *p);`

Input Arguments: pcb struct of process stack to be freed

Return Values: nothing

shell.h

#pragma once

#include <signal.h> // sigaction, sigemptyset, sigfillset, signal

#include <stdio.h> // dprintf, fputs, perror

#include <stdbool.h> // boolean

#include <stdlib.h> // malloc, free

#include <sys/time.h> // setitimer

#include <ucontext.h> // getcontext, makecontext, setcontext, swapcontext

#include <unistd.h> // read, usleep, write

#include <valgrind/valgrind.h>

#include <sys/wait.h>

#include <fcntl.h>

#include "kernel.h"

#include "dependencies.h"

#include "user_functions.h"

#include "parser.h"

#include "user.h"

#define INPUT_SIZE 4096 // maximum input buffer size

// process status definitions

#define STOPPED 3

#define RUNNING 2

#define SIG_TERMINATED 1

#define TERMINATED 0

// process foreground or background definitions

#define FG 0

#define BG 1

// define true or false

#define TRUE 1

#define FALSE 0

```
// defined the signals
```

```
#define S_SIGTERM 1
```

```
#define S_SIGSTOP 2
```

```
#define S_SIGCONT 3
```

```
extern int IS_BG; // global variable to state the status of the current active process
```

```
struct Job{
```

```
    int myPid; // job ID
```

```
    int JobNumber; // Counter for current job number since first job begins from 1
```

```
    int bgFlag; // FG = 0 and BG = 1
```

```
    struct Job *next; // pointer to next job
```

```
    char *commandInput; // argument
```

```
    int status; // tell whether its running or stopped
```

```
    int *pids; // list of pids of children
```

```
    int numChild; // number of children
```

```
    int *pids_finished; // boolean array list that checks every pid is finished
```

```
};
```

```
void setTimer(void);
```

Input Arguments: nothing

Return Values: nothing

```
void signalHandler(int signal);
```

Input Arguments: defined value of the signal

Return Values: nothing

```
void sigIntTermHandler(int signal);
```

Input Arguments: defined value of the signal

Return Values: nothing

`void sigcontHandler(int signal);`

Input Arguments: defined value of the signal

Return Values: nothing

`void sigtstpHandler(int signal);`

Input Arguments: defined value of the signal

Return Values: nothing

`void setSignalHandler(void);`

Input Arguments: nothing

Return Values: nothing

`void pennShredder(char* buffer);`

Input Arguments: input buffer

Return Values: nothing

`void pennShell();`

Input Arguments: nothing

Return Values: nothing

`struct Job *createJob(int pid, int bgFlag, int numChildren, char *input);`

Input Arguments: pid of job, bg or fg flag, number of children, and input array

Return Values: Job struct

`struct Job *addJob(struct Job *head, struct Job *newJob);`

Input Arguments: head of the jobs linked list in shell, job to be added

Return Values: Job struct of head of linked list in shell

`struct Job *removeJob(struct Job *head, int jobNum);`

Input Arguments: head of the jobs linked list in shell, job number of job to be removed

Return Values: Job struct of head of linked list in shell

```
struct Job *getJob(struct Job *head, int jobNum);
```

Input Arguments: head of the jobs linked list in shell, job number of job we are finding

Return Values: Job struct of job to be found

```
int getCurrentJob(struct Job *head);
```

Input Arguments: head of the jobs linked list in shell

Return Values: job number of chosen job to be returned

```
void changeStatus(struct Job *head, int jobNum, int newStatus);
```

Input Arguments: head of the jobs linked list in shell, job number of job whos status is to be changed, new status value

Return Values: Job struct of head of linked list in shell

```
void changeFGBG(struct Job *head, int jobNum, int newFGBG);
```

Input Arguments: head of the jobs linked list in shell, job number of job whos status is to be changed, new status value

Return Values: nothing

```
void iterateShell(struct Job *head);
```

Input Arguments: head of the jobs linked list in shell

Return Values: nothing

```
### user_functions.h  
#include <stdio.h>  
#include <stdlib.h>  
#include <string.h>  
#include <time.h>  
#include <dirent.h>  
#include <sys/stat.h>  
#include <time.h>  
#include <pwd.h>  
#include <grp.h>  
#include <ucontext.h>  
#include "scheduler.h"  
#include "dependencies.h"
```

```
void echoFunc(int argc, char *argv[]);
```

Input Arguments: number of arguments, array of strings of arguments

Return Values: nothing

```
void sleepFunc(int argc, char *argv[]);
```

Input Arguments: number of arguments, array of strings of arguments

Return Values: nothing

```
void busyFunc(void);
```

Input Arguments: nothing

Return Values: nothing

```
void idleFunc();
```

Input Arguments: nothing

Return Values: nothing

```
// ==== filesystem ====
```

```
void catFunc(int argc, char **argv);
```

Input Arguments: number of arguments, array of strings of arguments

Return Values: nothing

```
void lsFunc(int argc, char **argv);
```

Input Arguments: number of arguments, array of strings of arguments

Return Values: nothing

```
void touchFunc(int argc, char **argv);
```

Input Arguments: number of arguments, array of strings of arguments

Return Values: nothing

```
void mvFunc(int argc, char **argv);
```

Input Arguments: number of arguments, array of strings of arguments

Return Values: nothing

```
void cpFunc(int argc, char **argv);
```

Input Arguments: number of arguments, array of strings of arguments

Return Values: nothing

```
void rmFunc(int argc, char **argv);
```

Input Arguments: number of arguments, array of strings of arguments

Return Values: nothing

```
void chmodFunc(int argc, char **argv);
```

Input Arguments: number of arguments, array of strings of arguments

Return Values: nothing

```
void psFunc (int argc, char **argv);
```

Input Arguments: number of arguments, array of strings of arguments

Return Values: nothing

```
void killFunc (int argc, char **argv);
```

Input Arguments: number of arguments, array of strings of arguments

Return Values: nothing

```
void man();
```

Input Arguments: nothing

Return Values: nothing

```
void zombify(int argc, char **argv);
```

Input Arguments: number of arguments, array of strings of arguments

Return Values: nothing

```
void zombie_child();
```

Input Arguments: nothing

Return Values: nothing

```
void orphanify(int argc, char **argv);
```

Input Arguments: number of arguments, array of strings of arguments

Return Values: nothing

```
void orphan_child();
```

Input Arguments: nothing

Return Values: nothing

```
void niceFunc(char *argv[]);
```

Input Arguments: number of arguments, array of strings of arguments

Return Values: nothing

```
int nice_pid(char *argv[]);
```

Input Arguments: array of strings of arguments

Return Values: nothing

```
void logout();
```

Input Arguments: nothing

Return Values: nothing

user.h

#pragma once

#include <stdio.h>

#include <string.h>

#include <ucontext.h>

#include "scheduler.h"

#include "pcb.h"

#include "parser.h"

#include "dependencies.h"

#define MAX_CMD_LENGTH 1000 // maximum possible length of commands

#define MAX_ARGS 10 // maximum possible number of arguments

// definitions of the signals

#define S_SIGTERM 1

#define S_SIGSTOP 2

#define S_SIGCONT 3

// global variable for actively running process

extern Process *activeProcess;

#define PROMPT "\$ " // prompt definition

#define BUFFERSIZE 4096 // maximum size of buffer

char* concat(int argc, char *argv[]);

Input Arguments: number of arguments, array of arguments in form of strings

Return Values: returns concatenated string

pid_t p_spawn(void (*func)(), char *argv[], int fd0, int fd1);

Input Arguments: function to be run, array of arguments, input and output file descriptors

Return Values: pid of newly p_spawned process

```
pid_t p_waitpid(pid_t pid, int *wstatus, bool nohang);
```

Input Arguments: pid of process to be waited on, new status of process waited on, hang or no hanging wait

Return Values: pid of process that was waited on

```
void p_sleep(unsigned int ticks1);
```

Input Arguments: number of ticks left in running

Return Values: nothing

```
int p_kill(pid_t pid, int sig);
```

Input Arguments: pid of process to be killed, signal value in int

Return Values: -1 on failure, or return value of k_process_kill

```
void p_exit(void);
```

Input Arguments: nothing

Return Values: nothing

```
int p_nice(pid_t pid, int priority);
```

Input Arguments: pid of process priority to be changed, new priority

Return Values: returns 0 or -1 based on passing or failure

Filesystem Working Notes

File layout:

- src/
 - pennfat.c - entrypoint for the pennfat standalone
- fs/
 - fat.c/.h - kernel-level fat functions, file struct defs
 - user.c/.h - user-level fat functions

Namespacing

All kernel-level functions will start with the prefix `fs_`.

All user-level functions will start with the prefix `f_`.

Error Codes

PennOS specific filesystem errors will be in the 1000-2000 range.

```
```c
```

```
#define PEHOSTFS 1001 // could not open/close file in host filesystem
```

```
#define PEHOSTIO 1002 // could not perform I/O in host filesystem
```

```
#define PEBADFS 1003 // invalid PennFAT file, or was otherwise unable to mount
```

```
#define PENOFIL 1004 // specified file does not exist
```

```
#define PEINUSE 1005 // the specified file is in use in another context and an exclusive operation
was called
```

```
#define PETOOFAT 1006 // the filesystem is too fat and has no space for a new file
```

```
#define PEFMODE 1007 // attempted operation on a file in the wrong mode
```

```
#define PEFPERM 1008 // attempted operation on a file without read/write permissions
```

```
#define PEFNAME 1009 // the filename is invalid
```

```
#define PETOOMANYF 1101 // you have too many files open already
#define PESTDIO 1102 // tried to read from stdout or write to stdin
...
```

## Structs

When a file is opened, it returns a `struct file`:

```
```c
typedef struct file {
    filestat_t *entry; // mmaped to file entry in directory
    uint32_t offset; // current seek position
    int mode;
    uint8_t stdiomode; // 0 = FAT file, 1 = stdout, 2 = stdin
} file_t;
...
```

where `filestat_t` is the directory entry defined in the PennOS handout:

```
```c
typedef struct filestat {
 char name[32];
 uint32_t size;
 uint16_t blockno;
 uint8_t type;
 uint8_t perm;
 time_t mtime;
 uint8_t unused[16];
} filestat_t;
...
```

The OS should maintain a file descriptor table for each process linking an int to one of these `struct file`s`.

The low-level filesystem implementation operates using pointers to open file structs.

### ## stdin/stdout

Each `file_t *` struct contains information on whether it is a special file for reading from/writing to stdin/stdout.

The user-level functions `f_read()` and `f_write()` should check this information and redirect to the C API where

necessary rather than the FAT API.

Each process, on creation, should set entries `PSTDIN_FILENO` and `PSTDOUT_FILENO` in its PCB's fd table to file structs

with the correct flag set. This allows for later redirecting stdin/stdout by overwriting the entries in the fd table

with file structs linked to files on the FAT filesystem.

### ## File Locking Mechanism

In order to prevent multiple writers/conflicting read-writes, the PennFAT filesystem grants exclusive access to any

process that opens a file in a writing mode, and shared access to processes opening a file in read mode. To accomplish

this, the filesystem keeps a record of what files have been opened and in what mode; if a call to `fs_open()` would

violate the locking semantics, the syscall fails with an error.

This record is a `file_t *` array that utilizes array doubling to grow dynamically (initially sized at 4).

### ## Standalone

The standalone completes the demo plan with no "definitely/indirectly/possibly lost" memory leaks.

## ## Syscalls

### int f\_open(const char \*fname, int mode)

Open a file. If the file is opened in F\_WRITE or F\_APPEND mode, the file is created if it does not exist.

### \*\*Parameters\*\*

- `name`: the name of the file to open
- `mode`: the mode to open the file in (F\_WRITE, F\_READ, or F\_APPEND).

### \*\*Returns\*\*

the file descriptor of the opened file, -1 on error

### \*\*Exceptions\*\*

- `PENOFILE`: the requested file was in read mode and does not exist
- `PEHOSTIO`: failed to read from/write to host filesystem
- `PETOOFAT`: the operation would make a new file but the filesystem is full
- `PEFNAME`: the operation would make a new file but the filename is invalid
- `PEINUSE`: the requested file was opened in an exclusive mode and is currently in use

### int f\_close(int fd)

Closes the specified file, freeing any associated memory.

### \*\*Parameters\*\*

- `fd`: the file to close

### \*\*Returns\*\*

0 on a success, -1 on error.

### \*\*Exceptions\*\*

- `EINVAL`: the file descriptor is invalid

```
ssize_t f_read(int fd, int n, char *buf)
```

Read up to `n` bytes from the specified file into `buf`.

**\*\*Parameters\*\***

- `fd`: the file to read from
- `n`: the maximum number of bytes to read
- `buf`: a buffer to store the read bytes

**\*\*Returns\*\***

the number of bytes read; -1 on error

**\*\*Exceptions\*\***

- `EPEERM`: you do not have permission to read this file
- `EHOSTIO`: failed to read from host filesystem

```
ssize_t f_write(int fd, const char *str, ssize_t n)
```

Write up to `n` bytes from `buf` into the specified file.

**\*\*Parameters\*\***

- `fd`: the file to write to
- `str`: a buffer storing the bytes to write
- `b`: the maximum number of bytes to write

**\*\*Returns\*\***

the number of bytes written; -1 on error

**\*\*Exceptions\*\***

- `PEFMODE`: the file is not in write or append mode
- `PEFPERM`: you do not have permission to write to this file
- `PEHOSTIO`: failed to read from host filesystem
- `PETOOFAT`: filesystem is full

```
int f_unlink(const char *fname)
```

Removes the file with the given name.

**\*\*Parameters\*\***

- `fname`: the name of the file to delete

**\*\*Returns\*\***

0 on success; -1 on error

**\*\*Exceptions\*\***

- `PENOFIL`: the specified file does not exist
- `PEHOSTIO`: failed to i/o with the host entry

```
uint32_t f_lseek(int fd, int offset, int whence)
```

Seek the file offset to the given position, given by `offset`. If `whence` is `F\_SEEK\_SET` this is relative to the

start of the file, for `F\_SEEK\_CUR` relative to the current position, and for `F\_SEEK\_END` relative to the end of the

file.

**\*\*Parameters\*\***

- `fd`: the file to seek
- `offset`: where to seek to relative to `whence`
- `whence`: the seek mode



### **\*\*Returns\*\***

the new location in bytes from start of file; -1 on error

### **\*\*Exceptions\*\***

- `PEINVAL`: whence is not a valid option

```
filestat_t **f_ls(const char *fname)
```

Gets information for a file. If `fname` is NULL, gets information for all the files.

It is the caller's responsibility to free each of the returned structs. Use the convenience function `f\_frees()` to

do this quickly.

### **\*\*Parameters\*\***

- `fname`: the name of the file to get the stat of, or NULL to list all files

### **\*\*Returns\*\***

a pointer to an array of filestat struct pointers. The array will always be terminated with a NULL pointer.

### **\*\*Exceptions\*\***

- `PEHOSTIO`: failed to read from host filesystem

```
void f_frees(filestat_t **stat)
```

Free the filestat list returned by `f\_ls()`.

```
int f_rename(const char *oldname, const char *newname)
```

Rename a file.

### **\*\*Parameters\*\***

- `oldname`: the old name
- `newname`: the new name

#### **\*\*Returns\*\***

0 on success, -1 on failure

#### **\*\*Exceptions\*\***

- `PENOFILE`: the file does not exist
- `PEHOSTIO`: failed to perform IO on host drive
- `PEFNAME`: the new name is invalid

### int f\_chmod(int fd, char mode, uint8\_t bitset)

Edit the I/O permissions of an open file.

#### **\*\*Parameters\*\***

- `fd`: the file whose permissions to edit
- `mode`: the mode to edit it in; '+', '=', or '-'
- `bitset`: the permission bitset to edit by (a combination of FAT\_EXECUTE, FAT\_WRITE, and FAT\_READ)

#### **\*\*Returns\*\***

0 on success, -1 on error

#### **\*\*Exceptions\*\***

- `PEINVAL`: the mode is invalid

#### **### Kernel-Level Functions**

Kernel-level functions are documented inline in `src/fs/fat.c`.